




Sample benchmark: temperature comparison

—●— Main (T = 5000) —■— T = 1000 —▲— T = 3000 - - - Native reference



Mean $\pm$ 1 SD across 1,000 random subsets; n = 1-689. Lower is better. One best-scoring model retained per run. Logarithmic x-axis.